

# Oil\_Time\_Series.R

*rstudio-user*

*2019-12-23*

```
# Author:
#   RIHAD VATIAWA
#
# Time Series is a sequence of well-defined data points measured at consistent time
# intervals over a period of time. ... Time series analysis is the use of statistical
# methods to analyze time series data and extract meaningful statistics and characteristics about the d
#
# The goal of this project was to use various (ts) time series methods in order to
# make predictions on a dataset of our choice. The modeling techniques used were:
#
# ARIMA models
# Exponential Smoothing models
# Facebooks Prophet model
#
# The data sourced for our ts project came from the link below
# This data consists of monthly WTI (West Texas Intermediate) oil prices from Cushing,
# Oklahoma. The data extends from Jan 1986 upto 21 Dec 2019.
#   https://fred.stlouisfed.org/series/MCOILWTICO
#
# --- --- --- --- --- --- --- --- --- --- --- --- --- --- --- #

# import libraries
library(fpp2)

## Loading required package: ggplot2

## Loading required package: forecast

## Registered S3 method overwritten by 'xts':
##   method      from
##   as.zoo.xts zoo

## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo

## Registered S3 methods overwritten by 'forecast':
##   method      from
##   fitted.fracdiff   fracdiff
##   residuals.fracdiff fracdiff

## Loading required package: fma

## Loading required package: expsmooth
library(astsa)

##
## Attaching package: 'astsa'

## The following object is masked from 'package:fpp2':
```

```

##
## oil
## The following objects are masked from 'package:fma':
##
## chicken, sales
## The following object is masked from 'package:forecast':
##
## gas
library(tseries)
library(forecast)
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
## filter, lag
## The following objects are masked from 'package:base':
##
## intersect, setdiff, setequal, union
library(prophet)

## Loading required package: Rcpp
## Loading required package: rlang
# load in data
oil <- read.csv('MCOILWTICO.csv')

# preview the data
head(oil)


##          DATE MCOILWTICO
## 1 1986-01-01      22.93
## 2 1986-02-01      15.46
## 3 1986-03-01      12.61
## 4 1986-04-01      12.84
## 5 1986-05-01      15.38
## 6 1986-06-01      13.43

# glimpse(oil)

# transforming data into ts
oil_ts <- ts(oil$MCOILWTICO, start = c(1986, 1), frequency = 12)

# visualizing our ts data
plot(oil_ts)


```



```
# training data - all data except last six months: YR, MONTH
train <- window(oil_ts, c(1986, 1), c(2017, 8))
# plot(train)


# testing data - just the last six months: YR, MONTH
test <- window(oil_ts, c(2017, 9), c(2018, 2))
# plot(test)

# visualizing train and test data
plot(train)
lines(test, col = 'red')
```




```
#####  
# ARIMA Model  
#####
```

```
## Step 1: Visualize the ts  
# for obvious trend and increasing variance  
plot(oil_ts)
```



```
## Step 2: Check for stationarity  
# using a diff and a log to stationarize the data  
plot(diff(log(train)))
```



```
adf.test(train)
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: train  
## Dickey-Fuller = -2.2622, Lag order = 7, p-value = 0.4664  
## alternative hypothesis: stationary
```

```
# passes the Dickey Fuller test
```


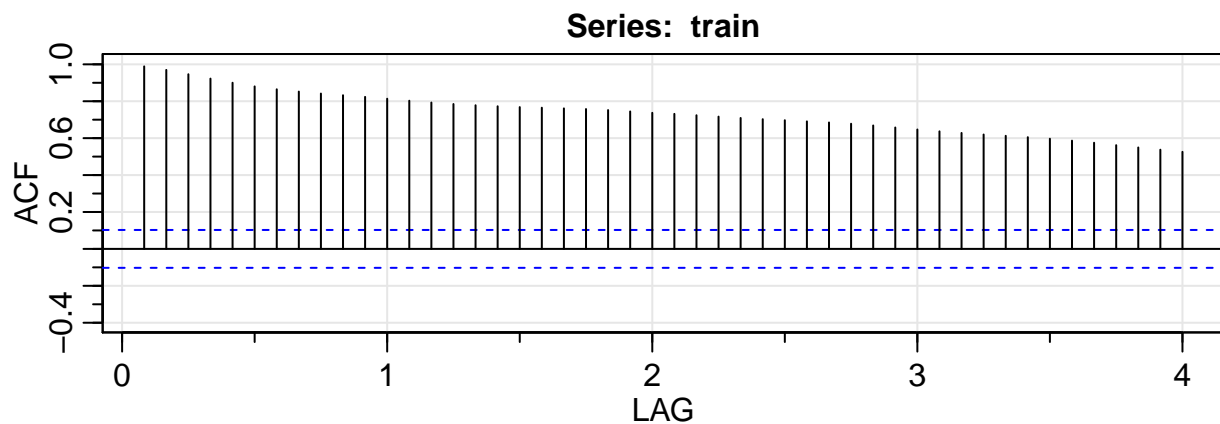
```
adf.test(diff(log(train)))
```

```
## Warning in adf.test(diff(log(train))): p-value smaller than printed p-value
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: diff(log(train))  
## Dickey-Fuller = -7.7017, Lag order = 7, p-value = 0.01  
## alternative hypothesis: stationary
```

```
# acf tails off - pacf kinda tails sorta
```

```
acf2(train)
```



```
##      ACF  PACF  
## [1,] 0.99  0.99  
## [2,] 0.97 -0.39  
## [3,] 0.95 -0.06  
## [4,] 0.92  0.08
```

```

## [5,] 0.90 0.03
## [6,] 0.88 0.07
## [7,] 0.86 0.12
## [8,] 0.85 0.01
## [9,] 0.84 0.00
## [10,] 0.83 0.04
## [11,] 0.82 -0.03
## [12,] 0.81 -0.03
## [13,] 0.80 0.02
## [14,] 0.79 0.11
## [15,] 0.78 0.04
## [16,] 0.78 0.03
## [17,] 0.77 0.01
## [18,] 0.77 0.01
## [19,] 0.77 0.04
## [20,] 0.76 -0.02
## [21,] 0.76 0.03
## [22,] 0.75 -0.08
## [23,] 0.74 0.00
## [24,] 0.74 0.08
## [25,] 0.73 0.09
## [26,] 0.72 -0.20
## [27,] 0.72 0.08
## [28,] 0.71 -0.01
## [29,] 0.70 0.03
## [30,] 0.70 0.04
## [31,] 0.69 -0.06
## [32,] 0.69 0.03
## [33,] 0.68 -0.06
## [34,] 0.67 -0.08
## [35,] 0.66 -0.01
## [36,] 0.65 0.05
## [37,] 0.64 0.03
## [38,] 0.63 0.03
## [39,] 0.62 -0.06
## [40,] 0.61 0.01
## [41,] 0.61 -0.08
## [42,] 0.60 -0.03
## [43,] 0.59 -0.04
## [44,] 0.58 -0.01
## [45,] 0.56 -0.03
## [46,] 0.55 0.05
## [47,] 0.54 -0.02
## [48,] 0.53 -0.05

```

#### *## ACF and PACF*

*# Autocorrelation refers to how correlated a ts (future value) is with its past values  
# whereas the ACF is the plot used to see the correlation (relationship) between  
# the points, upto and including the lag unit.  
# In ACF, the correlation coefficient is in the y-axis whereas the number  
# of lags is shown in the x-axis.*

*# suggests: p=2, d=1, q=2, P=0, D=0, Q=2, S=12*  
`auto.arima(train)`

```

## Series: train
## ARIMA(2,1,2)(0,0,2)[12]
##
## Coefficients:
##          ar1      ar2      ma1      ma2      sma1      sma2
##          1.5596 -0.6568 -1.1949  0.2760  0.1019 -0.1567
## s.e.    0.0829   0.0787   0.1039  0.0976  0.0561   0.0537
##
## sigma^2 estimated as 14.86:  log likelihood=-1046.74
## AIC=2107.49   AICc=2107.79   BIC=2135.05


# first model ----
fit <- sarima(log(train), p=2, d=1, q=2, P=0, D=0, Q=2, S=12)

## initial  value -2.469201
## iter    2 value -2.473196
## iter    3 value -2.507314
## iter    4 value -2.507671
## iter    5 value -2.508150
## iter    6 value -2.508465
## iter    7 value -2.509452
## iter    8 value -2.510230
## iter    9 value -2.510993
## iter   10 value -2.511360
## iter   11 value -2.511628
## iter   12 value -2.511816
## iter   13 value -2.511859
## iter   14 value -2.511885
## iter   15 value -2.511928
## iter   16 value -2.511982
## iter   17 value -2.511991
## iter   18 value -2.511996
## iter   19 value -2.511999
## iter   20 value -2.512001
## iter   21 value -2.512003
## iter   22 value -2.512004
## iter   23 value -2.512006
## iter   24 value -2.512007
## iter   25 value -2.512007
## iter   25 value -2.512007
## iter   25 value -2.512007
## final   value -2.512007
## converged
## initial  value -2.473238
## iter    2 value -2.473383
## iter    3 value -2.475430
## iter    4 value -2.475865
## iter    5 value -2.476714
## iter    6 value -2.478409
## iter    7 value -2.479678
## iter    8 value -2.479849
## iter    9 value -2.479890
## iter   10 value -2.479954
## iter   11 value -2.479958
## iter   12 value -2.479960

```

```
## iter 13 value -2.479960
## iter 14 value -2.479962
## iter 15 value -2.479966
## iter 16 value -2.479977
## iter 17 value -2.479981
## iter 18 value -2.479989
## iter 19 value -2.479990
## iter 20 value -2.479996
## iter 21 value -2.480006
## iter 22 value -2.480017
## iter 23 value -2.480025
## iter 24 value -2.480034
## iter 25 value -2.480054
## iter 26 value -2.480218
## iter 27 value -2.480314
## iter 28 value -2.480373
## iter 29 value -2.480417
## iter 30 value -2.480507
## iter 31 value -2.480630
## iter 32 value -2.480673
## iter 33 value -2.480689
## iter 34 value -2.480716
## iter 35 value -2.480801
## iter 36 value -2.481024
## iter 37 value -2.481400
## iter 38 value -2.482000
## iter 39 value -2.482356
## iter 40 value -2.483354
## iter 41 value -2.484759
## iter 42 value -2.485638
## iter 43 value -2.487146
## iter 44 value -2.487382
## iter 45 value -2.487570
## iter 46 value -2.487763
## iter 47 value -2.488081
## iter 48 value -2.488452
## iter 49 value -2.488586
## iter 50 value -2.488598
## iter 51 value -2.488609
## iter 52 value -2.488610
## iter 53 value -2.488610
## iter 54 value -2.488610
## iter 55 value -2.488611
## iter 56 value -2.488611
## iter 56 value -2.488611
## final value -2.488611
## converged
```





```
# preview AIC and BIC of model
print(cbind(fit$AIC, fit$BIC))
```


```
##           [,1]      [,2]
## [1,] -2.097129 -2.014015
```

```
fit$tttable
```

```
##           Estimate      SE t.value p.value
## ar1          1.2917 0.1657  7.7968 0.0000
## ar2         -0.3973 0.1621 -2.4502 0.0147
## ma1         -1.0174 0.1792 -5.6768 0.0000
## ma2          0.0921 0.1757  0.5244 0.6003
## sma1          0.0448 0.0577  0.7762 0.4381
## sma2         -0.0676 0.0545 -1.2398 0.2158
## constant     0.0024 0.0030  0.7891 0.4306
```

```
sarima_oil <- as.ts(sarima.for(train, n.ahead = 6, p=2, d=1, q=2, P=0, D=0, Q=2, S=12))
```

```
# compare predicted and actual
lines(test, col = 'blue')
```




```
# preview accuracy of model
accuracy(sarima_oil$pred, test)
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set 6.304671 7.111124 6.304671 10.63671 10.63671 0.4423583 2.029021
```

```
# second model ----
fit2 <- sarima(log(train), p=2, d=1, q=0, P=0, D=0, Q=0, S=12)
```

```
## initial value -2.469201
## iter 2 value -2.501140
## iter 3 value -2.504276
## iter 4 value -2.504299
## iter 5 value -2.504319
## iter 6 value -2.504330
## iter 7 value -2.504331
## iter 8 value -2.504332
## iter 8 value -2.504332
## final value -2.504332
## converged
## initial value -2.476313
## iter 2 value -2.476779
## iter 3 value -2.476819
## iter 4 value -2.476884
## iter 5 value -2.476896
## iter 6 value -2.476897
## iter 6 value -2.476897
## iter 6 value -2.476897
```

```
## final value -2.476897
## converged
```



```
# preview AIC and BIC of model
print(cbind(fit2$AIC, fit2$BIC))
```


```
##           [,1]      [,2]
## [1,] -2.094808 -2.053251
```

```
fit2$tttable
```

```
##      Estimate      SE t.value p.value
## ar1      0.2932 0.0524  5.5945  0.0000
## ar2     -0.0303 0.0528 -0.5744  0.5661
## constant  0.0016 0.0059  0.2773  0.7817
```

```
sarima_oil2 <- as.ts(sarima.for(train, n.ahead = 6, p=2, d=1, q=0, P=0, D=0, Q=1, S=12))
```

```
# compare predicted and actual
lines(test, col = 'blue')
```



```
# preview accuracy of model
accuracy(sarima_oil2$pred, test)
```

```
##               ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set 7.242352 8.530619 7.242352 12.0975 12.0975 0.5049333 2.414631
```

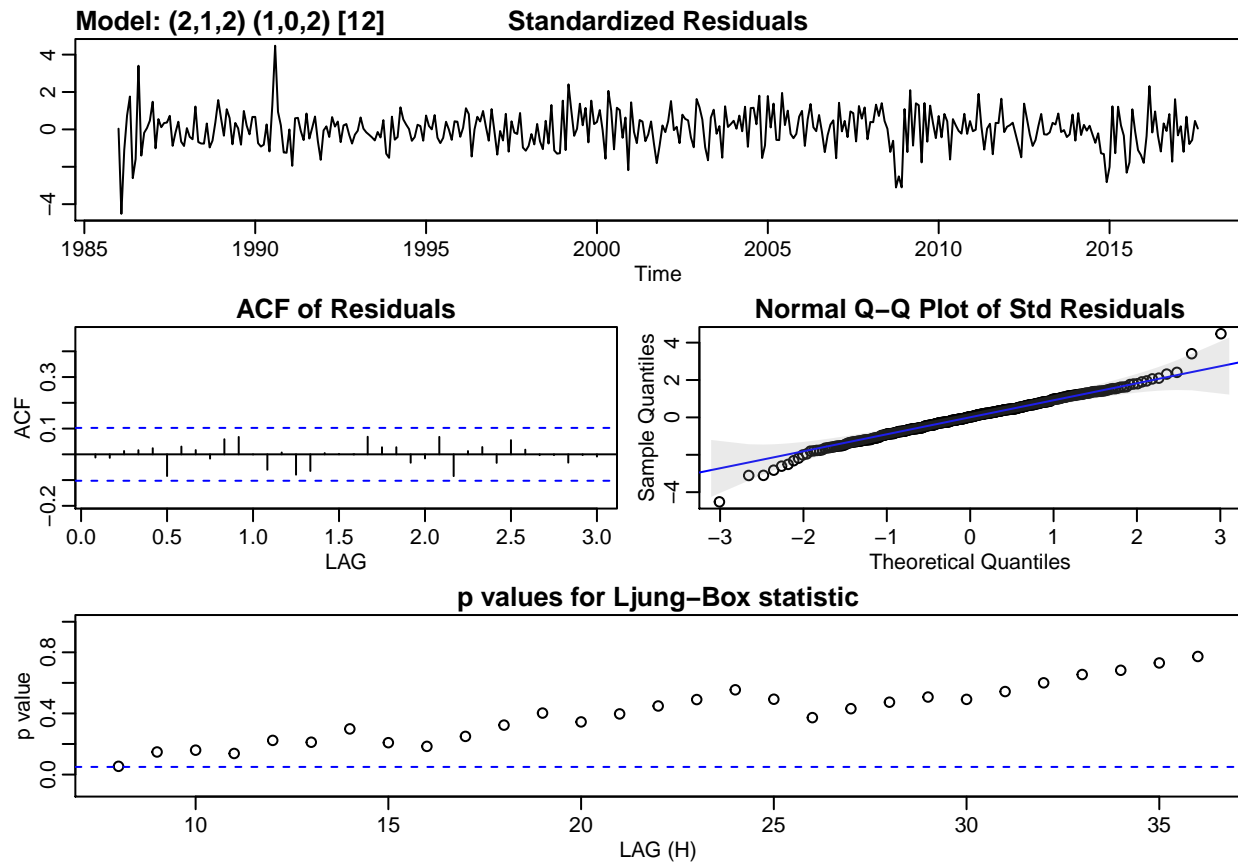
```
# third model ----
fit3 <- sarima(log(train), p=2, d=1, q=2, P=1, D=0, Q=2, S=12)
```

```
## initial value -2.485585
## iter 2 value -2.493045
## iter 3 value -2.537189
## iter 4 value -2.537724
## iter 5 value -2.538939
## iter 6 value -2.539211
## iter 7 value -2.541125
## iter 8 value -2.542098
## iter 9 value -2.542872
## iter 10 value -2.543224
## iter 11 value -2.543589
## iter 12 value -2.544106
## iter 13 value -2.544279
## iter 14 value -2.544345
## iter 15 value -2.544400
## iter 16 value -2.544445
## iter 17 value -2.544546
## iter 18 value -2.544610
## iter 19 value -2.544650
```

```
## iter 20 value -2.544763
## iter 21 value -2.544814
## iter 22 value -2.544848
## iter 23 value -2.544934
## iter 24 value -2.545387
## iter 25 value -2.545919
## iter 26 value -2.546460
## iter 27 value -2.547245
## iter 28 value -2.547501
## iter 29 value -2.547810
## iter 30 value -2.549027
## iter 31 value -2.549815
## iter 32 value -2.550803
## iter 33 value -2.551083
## iter 34 value -2.551307
## iter 35 value -2.551353
## iter 36 value -2.551371
## iter 37 value -2.551409
## iter 38 value -2.551456
## iter 39 value -2.551512
## iter 40 value -2.551537
## iter 41 value -2.551537
## iter 42 value -2.551538
## iter 43 value -2.551538
## iter 44 value -2.551538
## iter 45 value -2.551539
## iter 46 value -2.551539
## iter 47 value -2.551541
## iter 48 value -2.551548
## iter 49 value -2.551557
## iter 50 value -2.551639
## iter 51 value -2.551694
## iter 52 value -2.551711
## iter 53 value -2.551718
## iter 54 value -2.551725
## iter 55 value -2.551728
## iter 56 value -2.551730
## iter 57 value -2.551731
## iter 58 value -2.551732
## iter 59 value -2.551732
## iter 60 value -2.551733
## iter 61 value -2.551733
## iter 62 value -2.551733
## iter 63 value -2.551733
## iter 64 value -2.551733
## iter 65 value -2.551733
## iter 65 value -2.551733
## iter 65 value -2.551733
## final value -2.551733
## converged
## initial value -2.483153
## iter 2 value -2.484372
## iter 3 value -2.484456
## iter 4 value -2.485151
```

```
## iter    5 value -2.485515
## iter    6 value -2.485607
## iter    7 value -2.486053
## iter    8 value -2.486894
## iter    9 value -2.487590
## iter   10 value -2.488060
## iter   11 value -2.488326
## iter   12 value -2.488758
## iter   13 value -2.488780
## iter   14 value -2.488805
## iter   15 value -2.488814
## iter   16 value -2.488817
## iter   17 value -2.488820
## iter   18 value -2.488825
## iter   19 value -2.488825
## iter   20 value -2.488827
## iter   21 value -2.488828
## iter   22 value -2.488831
## iter   23 value -2.488835
## iter   24 value -2.488841
## iter   25 value -2.488844
## iter   26 value -2.488844
## iter   27 value -2.488845
## iter   28 value -2.488845
## iter   29 value -2.488846
## iter   30 value -2.488847
## iter   31 value -2.488848
## iter   32 value -2.488849
## iter   33 value -2.488849
## iter   34 value -2.488850
## iter   35 value -2.488852
## iter   36 value -2.488853
## iter   37 value -2.488854
## iter   38 value -2.488855
## iter   39 value -2.488855
## iter   40 value -2.488857
## iter   41 value -2.488860
## iter   42 value -2.488863
## iter   43 value -2.488864
## iter   44 value -2.488865
## iter   45 value -2.488866
## iter   46 value -2.488869
## iter   47 value -2.488877
## iter   48 value -2.488896
## iter   49 value -2.488905
## iter   50 value -2.488935
## iter   51 value -2.488953
## iter   52 value -2.488997
## iter   53 value -2.489051
## iter   54 value -2.489094
## iter   55 value -2.489208
## iter   56 value -2.489228
## iter   57 value -2.489424
## iter   58 value -2.489537
```

```
## iter 59 value -2.489696
## iter 60 value -2.489766
## iter 61 value -2.489854
## iter 62 value -2.490028
## iter 63 value -2.490299
## iter 64 value -2.490806
## iter 65 value -2.490868
## iter 66 value -2.491072
## iter 67 value -2.491154
## iter 68 value -2.491206
## iter 69 value -2.491244
## iter 70 value -2.491246
## iter 71 value -2.491250
## iter 72 value -2.491301
## iter 73 value -2.491392
## iter 74 value -2.491439
## iter 75 value -2.491458
## iter 76 value -2.491463
## iter 77 value -2.491466
## iter 78 value -2.491471
## iter 79 value -2.491477
## iter 80 value -2.491483
## iter 81 value -2.491485
## iter 82 value -2.491485
## iter 83 value -2.491486
## iter 84 value -2.491486
## iter 85 value -2.491487
## iter 86 value -2.491489
## iter 87 value -2.491489
## iter 88 value -2.491490
## iter 89 value -2.491490
## iter 90 value -2.491490
## iter 91 value -2.491490
## iter 92 value -2.491491
## iter 93 value -2.491492
## iter 94 value -2.491492
## iter 94 value -2.491492
## iter 94 value -2.491492
## final value -2.491492
## converged
```



```
# preview AIC and BIC of model
print(cbind(fit3$AIC, fit3$BIC))
```

```
##           [,1]      [,2]
## [1,] -2.097613 -2.00411
```


```
fit3$tttable
```

```
##           Estimate      SE t.value p.value
## ar1           1.2800 0.1658  7.7199 0.0000
## ar2          -0.3836 0.1628 -2.3558 0.0190
## ma1          -1.0034 0.1785 -5.6220 0.0000
## ma2           0.0755 0.1757  0.4297 0.6676
## sar1          -0.8884 0.1121 -7.9290 0.0000
## sma1           0.9457 0.1370  6.9043 0.0000
## sma2          -0.0067 0.0696 -0.0963 0.9234
## constant      0.0023 0.0031  0.7465 0.4559
```

```
sarima_oil3 <- as.ts(sarima.for(train, n.ahead = 6, p=2, d=1, q=2, P=1, D=0, Q=2, S=12))
```

```
# compare predicted and actual
lines(test, col = 'blue')
```





```
# preview accuracy of model
accuracy(sarima_oil3$pred, test)
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set 5.389568 6.054024 5.389568 9.107386 9.107386 0.3763519 1.730619
```

```
# forth model ----
fit4 <- sarima(log(train), p=2, d=1, q=1, P=1, D=0, Q=1, S=12)
```


```
## initial value -2.485585
## iter 2 value -2.487772
## iter 3 value -2.531278
## iter 4 value -2.532389
## iter 5 value -2.532546
## iter 6 value -2.532822
## iter 7 value -2.533527
## iter 8 value -2.534885
## iter 9 value -2.535602
## iter 10 value -2.536139
## iter 11 value -2.536745
## iter 12 value -2.537083
## iter 13 value -2.537138
## iter 14 value -2.537181
## iter 15 value -2.537191
## iter 16 value -2.537218
## iter 17 value -2.537221
## iter 18 value -2.537222
## iter 19 value -2.537222
```

```
## iter 20 value -2.537223
## iter 21 value -2.537226
## iter 22 value -2.537227
## iter 23 value -2.537228
## iter 24 value -2.537230
## iter 25 value -2.537235
## iter 26 value -2.537245
## iter 27 value -2.537262
## iter 28 value -2.537264
## iter 29 value -2.537280
## iter 30 value -2.537281
## iter 31 value -2.537289
## iter 32 value -2.537290
## iter 33 value -2.537290
## iter 34 value -2.537290
## iter 35 value -2.537290
## iter 36 value -2.537291
## iter 37 value -2.537292
## iter 38 value -2.537294
## iter 39 value -2.537297
## iter 40 value -2.537300
## iter 41 value -2.537301
## iter 42 value -2.537302
## iter 42 value -2.537302
## iter 42 value -2.537302
## final value -2.537302
## converged
## initial value -2.476130
## iter 2 value -2.477311
## iter 3 value -2.478376
## iter 4 value -2.478395
## iter 5 value -2.478416
## iter 6 value -2.478435
## iter 7 value -2.478467
## iter 8 value -2.478489
## iter 9 value -2.478512
## iter 10 value -2.478548
## iter 11 value -2.478634
## iter 12 value -2.478755
## iter 13 value -2.479018
## iter 14 value -2.479514
## iter 15 value -2.479611
## iter 16 value -2.479850
## iter 17 value -2.479885
## iter 18 value -2.480058
## iter 19 value -2.480196
## iter 20 value -2.480402
## iter 21 value -2.480672
## iter 22 value -2.481116
## iter 23 value -2.481888
## iter 24 value -2.482037
## iter 25 value -2.482118
## iter 26 value -2.482164
## iter 27 value -2.482172
```

```
## iter 28 value -2.482173
## iter 29 value -2.482177
## iter 30 value -2.482177
## iter 31 value -2.482178
## iter 32 value -2.482180
## iter 33 value -2.482181
## iter 34 value -2.482182
## iter 35 value -2.482184
## iter 36 value -2.482188
## iter 37 value -2.482191
## iter 38 value -2.482192
## iter 39 value -2.482192
## iter 40 value -2.482192
## iter 40 value -2.482193
## final value -2.482193
## converged
```

```
## Warning in sqrt(diag(fitit$var.coef)): NaNs produced
```

```
## Warning in sqrt(diag(fitit$var.coef)): NaNs produced
```



```
# preview AIC and BIC of model
print(cbind(fit4$AIC, fit4$BIC))
```


```
##           [,1]      [,2]
## [1,] -2.089569 -2.016844
```

```
fit4$table
```

```
##           Estimate      SE t.value p.value
## ar1        -0.1896    NaN      NaN      NaN
## ar2         0.1138    NaN      NaN      NaN
## ma1         0.4823    NaN      NaN      NaN
## sar1        -0.9049  0.0920 -9.8344  0.0000
## sma1         0.9661  0.0909 10.6311  0.0000
## constant    0.0015  0.0061  0.2444  0.8071
```

```
sarima_oil4 <- as.ts(sarima.for(train, n.ahead = 6, p=2, d=1, q=1, P=1, D=0, Q=1, S=12))
```

```
# compare predicted and actual
lines(test, col = 'blue')
```



```
# preview accuracy of model
accuracy(sarima_oil4$pred, test)
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set 9.655375 11.08977 9.655375 16.21246 16.21246 0.5282617 3.146215
```

```
# ARIMA Summary:
```

```
# first model was best - slightly less accurate than third model but less complex
```

```
# AR only model was not as good as first model even though suggested by ACF and PACF plots
```

```
#####
```

```
# Exponential Smoothing Model
```

```
#####
```

```

# now trying simple exponential smoothing due to trend in data

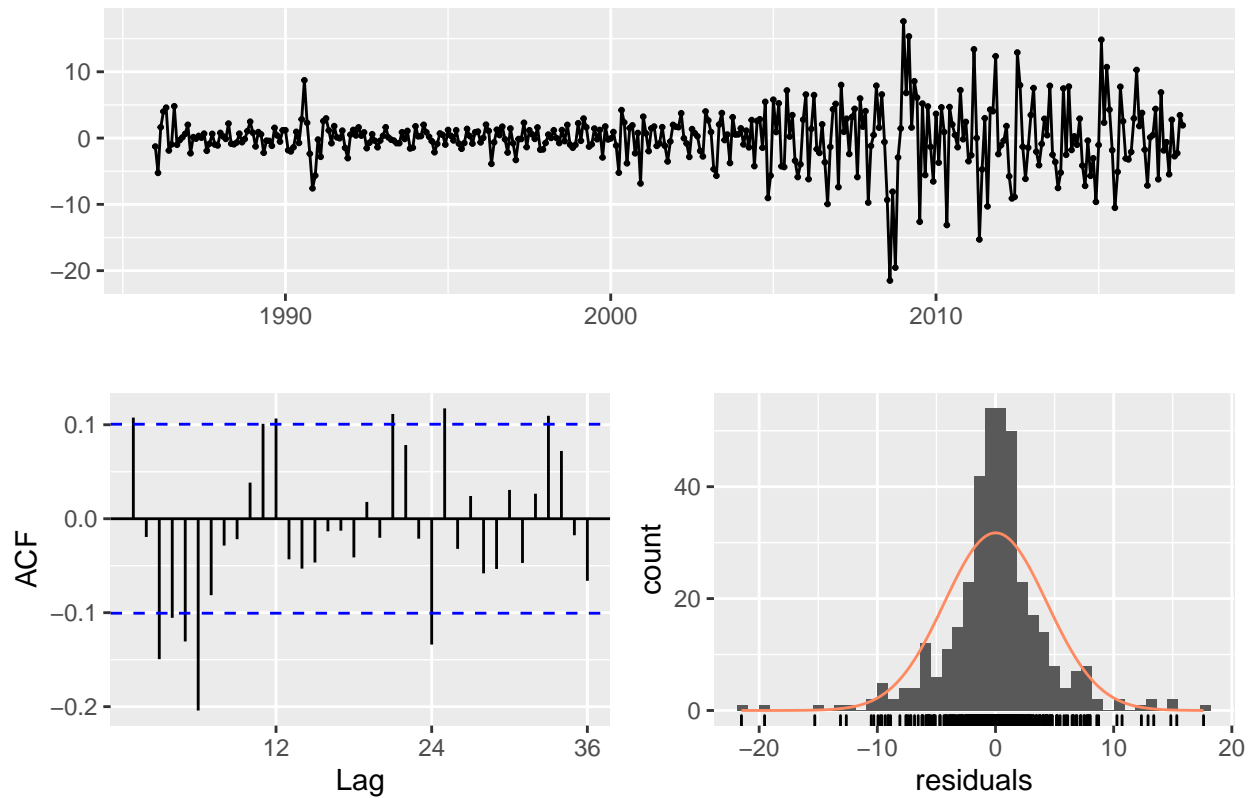
# holt model
fit_holt <- holt(train, h = 12)
summary(fit_holt)

##
## Forecast method: Holt's method
##
## Model Information:
## Holt's method
##
## Call:
## holt(y = train, h = 12)
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 0.4311
##
## Initial states:
##   l = 25.8715
##   b = -1.6742
##
## sigma: 4.3603
##
##      AIC      AICc      BIC
## 3382.374 3382.534 3402.075
##
## Error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.01214228 4.337284 2.891557 0.3268413 7.081585 0.2619254
##              ACF1
## Training set 0.1076263
##
## Forecasts:
##      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Sep 2017      48.35465 42.766707 53.94259 39.808629 56.90067
## Oct 2017      48.66949 38.914188 58.42479 33.750045 63.58893
## Nov 2017      48.98433 34.721287 63.24737 27.170888 70.79777
## Dec 2017      49.29917 30.125325 68.47301 19.975305 78.62303
## Jan 2018      49.61401 25.131823 74.09620 12.171736 87.05628
## Feb 2018      49.92885 19.758275 80.09943  3.786938 96.07076
## Mar 2018      50.24369 14.023525 86.46386 -5.150273 105.63765
## Apr 2018      50.55853  7.945226 93.17184 -14.612895 115.72996
## May 2018      50.87337  1.539313 100.20743 -24.576559 126.32330
## Jun 2018      51.18821 -5.179978 107.55640 -35.019495 137.39592
## Jul 2018      51.50305 -12.199931 115.20604 -45.922253 148.92836
## Aug 2018      51.81789 -19.509148 123.14493 -57.267403 160.90319

# checking for autocorrelation in residuals
checkresiduals(fit_holt)


```

## Residuals from Holt's method



```
##  
##  Ljung-Box test  
##  
## data:  Residuals from Holt's method  
## Q* = 71.252, df = 20, p-value = 1.136e-07  
##  
## Model df: 4.    Total lags used: 24  
# plotting model fit  
autoplot(fit_holt) + autolayer(fitted(fit_holt))
```

## Forecasts from Holt's method



```
# checking model accuracy
```

```
accuracy(fit_holt, test)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.01214228 4.337284 2.891557 0.3268413 7.081585 0.2619254
## Test set     7.83325097 9.060829 7.833251 13.1343054 13.134305 0.7095581
##              ACF1 Theil's U
## Training set 0.1076263      NA
## Test set     0.5199591 2.566832
```

```
# forecast test values
```


```
for_fit <- forecast(fit_holt, h = 6)
```

```
# plot the predicted data
```

```
plot(for_fit)
```

```
lines(test, col = 'red')
```

## Forecasts from Holt's method



```
# holt model damped
fit_holt_d <- holt(train, h = 12, damped = TRUE)
summary(fit_holt_d)
```

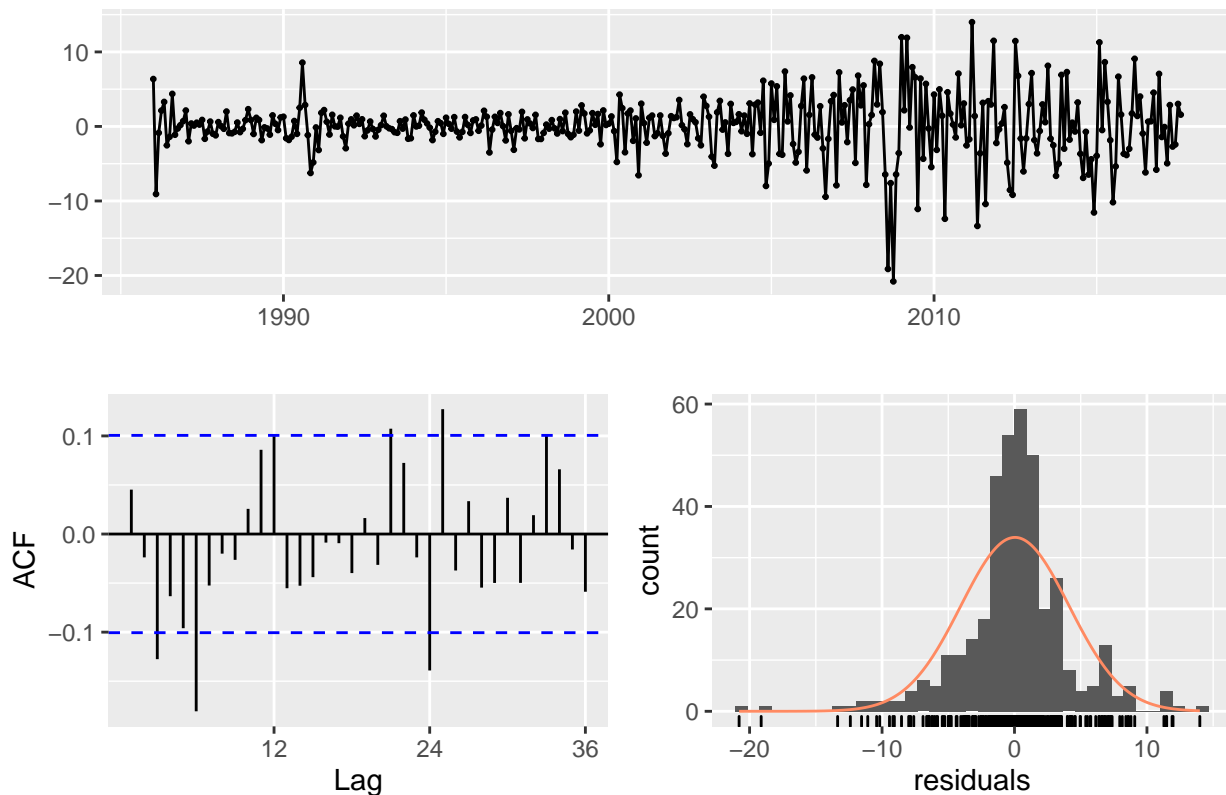
```
##
## Forecast method: Damped Holt's method
##
## Model Information:
## Damped Holt's method
##
## Call:
## holt(y = train, h = 12, damped = TRUE)
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 0.4508
##   phi   = 0.8
##
## Initial states:
##   l = 17.4339
##   b = -1.0832
##
## sigma: 4.1087
##
##      AIC      AICc      BIC
## 3338.185 3338.411 3361.826
##
## Error measures:
```



```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.03484041 4.081539 2.76284 0.1178085 6.779181 0.2502659
##           ACF1
## Training set 0.04524763
##
## Forecasts:
##           Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Sep 2017      48.47409 43.20863 53.73955 40.421268 56.52692
## Oct 2017      48.82149 39.93079 57.71220 35.224330 62.41865
## Nov 2017      49.09941 36.67217 61.52666 30.093582 68.10524
## Dec 2017      49.32175 33.43437 65.20913 25.024101 73.61940
## Jan 2018      49.49962 30.24936 68.74987 20.058894 78.94034
## Feb 2018      49.64191 27.14192 72.14191 15.231139 84.05269
## Mar 2018      49.75575 24.12745 75.38404 10.560652 88.95084
## Apr 2018      49.84682 21.21410 78.47954  6.056851 93.63678
## May 2018      49.91967 18.40486 81.43448  1.721931 98.11741
## Jun 2018      49.97796 15.69933 84.25658 -2.446674 102.40259
## Jul 2018      50.02458 13.09488 86.95429 -6.454529 106.50370
## Aug 2018      50.06189 10.58750 89.53627 -10.308970 110.43274
```

```
checkresiduals(fit_holt_d)
```


Residuals from Damped Holt's method



```
##
## Ljung-Box test
##
## data: Residuals from Damped Holt's method
## Q* = 52.815, df = 19, p-value = 4.991e-05
##
```

```
## Model df: 5.    Total lags used: 24
```

```
autoplot(fit_holt_d) + autolayer(fitted(fit_holt_d))
```




```
accuracy(fit_holt_d, test)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.03484041 4.081539 2.762840  0.1178085  6.779181  0.2502659
## Test set     7.83195437 9.125384 7.831954 13.1116442 13.111644  0.7094406
##              ACF1 Theil's U
## Training set 0.04524763      NA
## Test set     0.52498746  2.58097
```

```
for_fit2 <- forecast(fit_holt_d, h = 6)
plot(for_fit2)
lines(test, col = 'red')
```

## Forecasts from Damped Holt's method



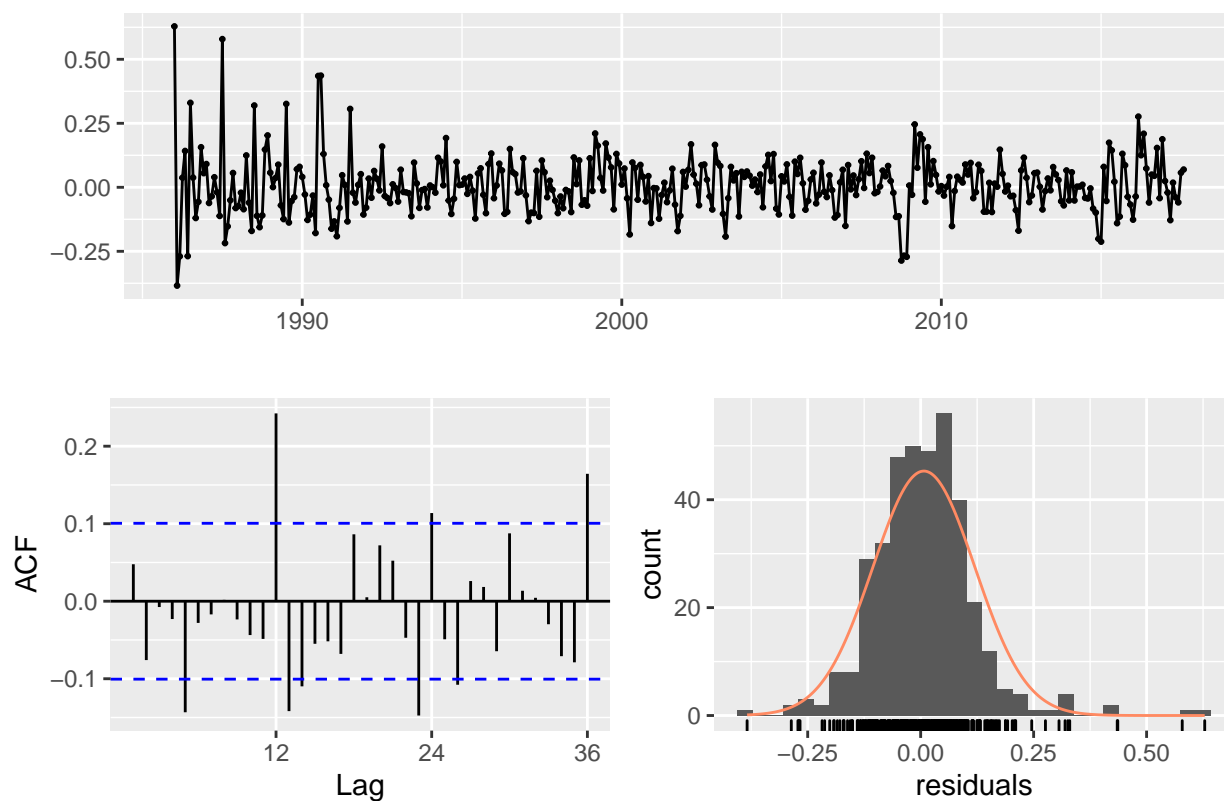
```
# holt winters model - dont expect this model to work super well - no clear seasonality in the data
fit_hw <- hw(train, seasonal = "multiplicative")
summary(fit_hw)
```

```
##
## Forecast method: Holt-Winters' multiplicative method
##
## Model Information:
## Holt-Winters' multiplicative method
##
## Call:
## hw(y = train, seasonal = "multiplicative")
##
## Smoothing parameters:
##   alpha = 0.9078
##   beta  = 0.0442
##   gamma = 0.0922
##
## Initial states:
##   l = 14.7024
##   b = 0.3698
##   s = 0.9437 0.9447 1.0892 1.0537 0.956 0.7575
##       1.2178 1.027 0.9909 1.0592 1.0262 0.9341
##
## sigma: 0.1155
##
##      AIC      AICc      BIC
## 3321.590 3323.281 3388.573
```

```
##
## Error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.09302371 4.623187 3.118393 -0.4995332 8.134427 0.2824729
##           ACF1
## Training set 0.3277393
##
## Forecasts:
##           Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Sep 2017      46.87253 39.9357465 53.80931 36.2636366 57.48142
## Oct 2017      45.79068 36.3602478 55.22111 31.3680780 60.21328
## Nov 2017      43.64089 32.4648272 54.81696 26.5485744 60.73322
## Dec 2017      42.67346 29.7537946 55.59313 22.9145357 62.43239
## Jan 2018      41.82564 27.2811861 56.37010 19.5818156 64.06947
## Feb 2018      43.10529 26.2080083 60.00258 17.2631264 68.94746
## Mar 2018      44.80010 25.2616627 64.33853 14.9186398 74.68156
## Apr 2018      45.58369 23.6801382 67.48725 12.0850960 79.08229
## May 2018      44.88387 21.2993539 68.46838 8.8144667 80.95327
## Jun 2018      44.28739 18.9912415 69.58354 5.6002680 82.97452
## Jul 2018      43.69738 16.6978727 70.69688 2.4051990 84.98955
## Aug 2018      42.45633 14.1934195 70.71925 -0.7680637 85.68073
## Sep 2018      41.43828 11.6052633 71.27130 -4.1873830 87.06395
## Oct 2018      40.42847 9.2939083 71.56303 -7.1877336 88.04468
## Nov 2018      38.47853 6.8552611 70.10179 -9.8850849 86.84214
## Dec 2018      37.57376 4.6880179 70.45950 -12.7206408 87.86815
## Jan 2019      36.77547 2.5597967 70.99115 -15.5528873 89.10383
## Feb 2019      37.84614 0.4753869 75.21689 -19.3074952 94.99978
## Mar 2019      39.27638 -1.8254895 80.37824 -23.5835025 102.13625
## Apr 2019      39.90331 -4.2945848 84.10121 -27.6915396 107.49817
## May 2019      39.23032 -6.7082190 85.16885 -31.0266112 109.48725
## Jun 2019      38.64812 -9.1484632 86.44470 -34.4504441 111.74668
## Jul 2019      38.07188 -11.6080362 87.75179 -37.9069928 114.05074
## Aug 2019      36.92968 -13.8739264 87.73329 -40.7677315 114.62709
```


```
checkresiduals(fit_hw)
```

## Residuals from Holt–Winters' multiplicative method



```
##
##  Ljung-Box test
##
## data:  Residuals from Holt-Winters' multiplicative method
## Q* = 74.791, df = 8, p-value = 5.432e-13
##
## Model df: 16.    Total lags used: 24
autoplot(fit_hw) + autolayer(fitted(fit_hw))
```

## Forecasts from Holt–Winters' multiplicative method



```
accuracy(fit_hw, test)
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set -0.09302371  4.623187  3.118393 -0.4995332  8.134427
## Test set     12.99024979 14.643891 12.990250 21.9058454 21.905845
##              MASE      ACF1 Theil's U
## Training set 0.2824729 0.3277393      NA
## Test set     1.1766936 0.5341790  4.167851
```

```
for_fit3 <- forecast(fit_hw, h = 6)
plot(for_fit3)
lines(test, col = 'red')
```

## Forecasts from Holt–Winters' multiplicative method



```
# Exponential Smoothing Summary:
# holt model seems to predict the best
# makes sense as there is a trend but no seasonality
```

```
#####
# Phophet Model
#####
```

```
head(oil)
```

```
##          DATE MCOILWTICO
## 1 1986-01-01      22.93
## 2 1986-02-01      15.46
## 3 1986-03-01      12.61
## 4 1986-04-01      12.84
## 5 1986-05-01      15.38
## 6 1986-06-01      13.43
```

```
# renaming columns for prophet modeling
```

```
colnames(oil)[1] <- 'ds'
colnames(oil)[2] <- 'y'
```

```
# re-subsetting data - prophet requires dataframe instead of ts
```

```
oil_train = oil[1:380,]
oil_test = oil[381:386,]
```

```
prophet_fit <- prophet(oil_train)
```

```
## Disabling weekly seasonality. Run prophet with weekly.seasonality=TRUE to override this.
```

```
## Disabling daily seasonality. Run prophet with daily.seasonality=TRUE to override this.
```

```
# make df for forecasted predictions
```


```
future <- make_future_dataframe(prophet_fit, periods = 6, freq = 'month')
```

```
forecast <- predict(prophet_fit, future)
```

```
tail(forecast[c('ds', 'yhat', 'yhat_lower', 'yhat_upper')])
```


```
##           ds      yhat yhat_lower yhat_upper
## 381 2017-09-01 55.63151  40.18352  71.22727
## 382 2017-10-01 54.09709  38.54682  70.55393
## 383 2017-11-01 51.38124  35.24158  66.29676
## 384 2017-12-01 49.72739  34.28923  65.42814
## 385 2018-01-01 50.31957  34.02240  66.20291
## 386 2018-02-01 49.30186  34.52141  65.95236
```

```
plot(prophet_fit, forecast)
```



```
prophet_plot_components(prophet_fit, forecast)
```





```
# checking model accuracy
accuracy(forecast[c('yhat')][380:386,], test)
```

```
##           ME      RMSE      MAE      MPE      MAPE      ACF1 Theil's U
## Test set 4.101113 8.577843 7.540457 6.059809 12.87085 0.5703779 2.291049
```

```
# Prophet model summary:
# fits better than exponential smoothing models but not as well as final ARIMA model
```